

OPENSEES-ON-DESIGNSAFE TRAINING

LEVERAGING TACC'S HPC RESOURCES

Silvia Mazzoni, PhD

Applications Specialist @ **DesignSafe**

18 September 2025



NSF NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

WORKSHOP OBJECTIVE

- Provide an overview of OpenSees on DesignSafe & TACC
- Introduce the new OpenSees-on-DesignSafe Training Document
- For me to gain a better understanding of how users approach and use OpenSees on DesignSafe

.... I hope you “walk” away with more questions than answers!



NSF NHERI
DESIGNSAFE



UCLA

TACC

RICE

Florida Tech

WHAT IS YOUR EXPERIENCE LEVEL

- OpenSees Tcl
- OpenSeesMP
- OpenSeesSP
- OpenSeesPy
- OpenSeesPy-MPI
- OpenSeesPy + concurrent futures
- OpenSees on DesignSafe Web Portal
- OpenSees on DS JupyterHub
- OpenSees & Tapis
- OpenSees on HPC via SSH & slurm Jobs

Low
Medium
High



NSF NHERI DESIGNSAFE



UCLA

TACC

RICE

Florida Tech

OUTLINE

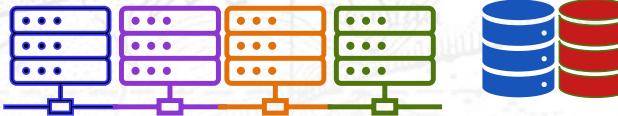
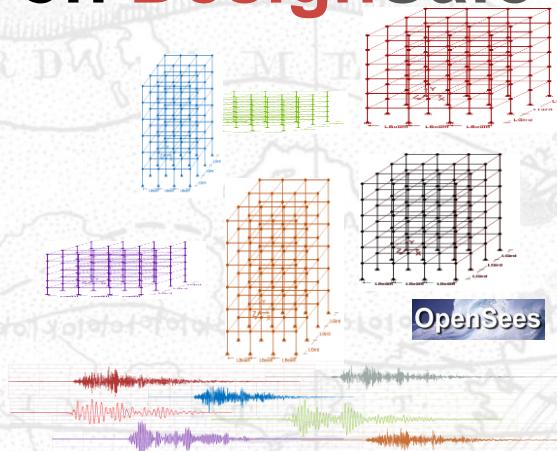
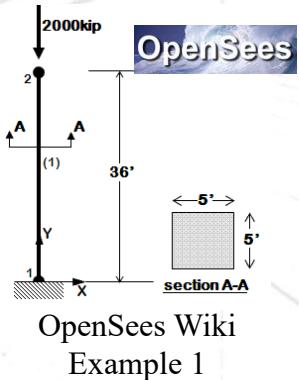
Why Run OpenSees on DesignSafe

When To Run OpenSees on DesignSafe

How To Run OpenSees on DesignSafe



Why Run OpenSees on DesignSafe



Scalable Resources & Environments



NSF NHERI DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

When To Run OpenSees on DesignSafe

Analysis Type	Memory Usage	CPU vs GPU	Parallelism Needs	Speed Sensitivity	Description + Reason / Notes
Parallel Monte Carlo	♦ Low–Moderate	<input checked="" type="checkbox"/> CPU preferred	♦ High (trivially parallel)	♦ Low–Moderate	Monte Carlo methods involve running repeated random samples to solve deterministic problems that are too complex for direct computation. Each run is independent; minimal memory overhead, great for CPU clusters or grid computing.
Parametric Sweeps	♦ Low–Moderate	<input checked="" type="checkbox"/> CPU preferred	♦ High	♦ Low–Moderate	Run a model or simulation across a range of input parameters to study system behavior or performance variation. Like Monte Carlo, easily parallelized with little inter-process communication.
ML Training Loops	♦ Moderate–High	<input type="radio"/> GPU accelerated	♦ Moderate–High	<input type="radio"/> High	ML training involves iterative updates to a model's parameters using a training dataset. GPU-accelerated for fast matrix ops; training requires high compute, memory usage depends on model size.
Coupled Simulations	<input type="radio"/> High	<input checked="" type="checkbox"/> CPU dominated	♦ Low–Moderate	<input type="radio"/> High	Simultaneously simulate multiple interacting physical domains or solvers (e.g., fluid + structure, heat + stress). Memory-intensive due to mesh/data exchange between solvers; limited parallelization unless domain-decomposed.
Stepwise Simulation	♦ Moderate	<input checked="" type="checkbox"/> CPU preferred	♦ Moderate	♦ Moderate	A sequential, time-marching simulation solving physical phenomena across a time domain, e.g., using finite difference, finite element, or finite volume methods. Each step may use iterative solvers; memory builds with history data.
Batch Pre/Post-Processing	♦ Low	<input checked="" type="checkbox"/> <input type="radio"/> CPU or GPU capable	♦ High	♦ Low–Moderate	Transform or clean a large batch of data files or inputs to prepare for further analysis or training. Lightweight tasks like data cleaning, often highly parallel and not memory-intensive.



NSF NHERI DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

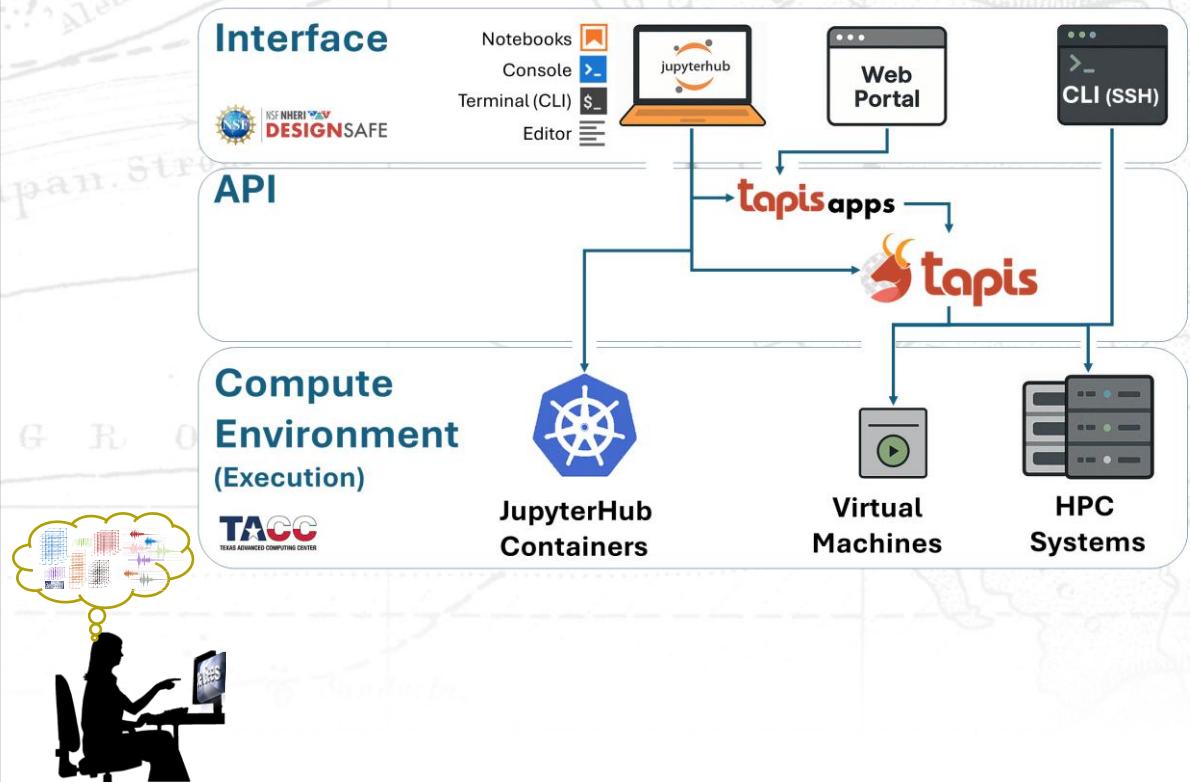
TACC

RICE

Florida Tech

Scalable & Adaptable Resources & Environments

How To Run OpenSees on DesignSafe



Interface Environments:

1. Jupyter Hub
2. Web Portal
3. CLI (SSH)

API:

1. Tapis Apps
2. Tapis

Compute Environments:

1. Jupyter Hub
2. Virtual Machines
3. HPC Systems on TACC



NSF NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

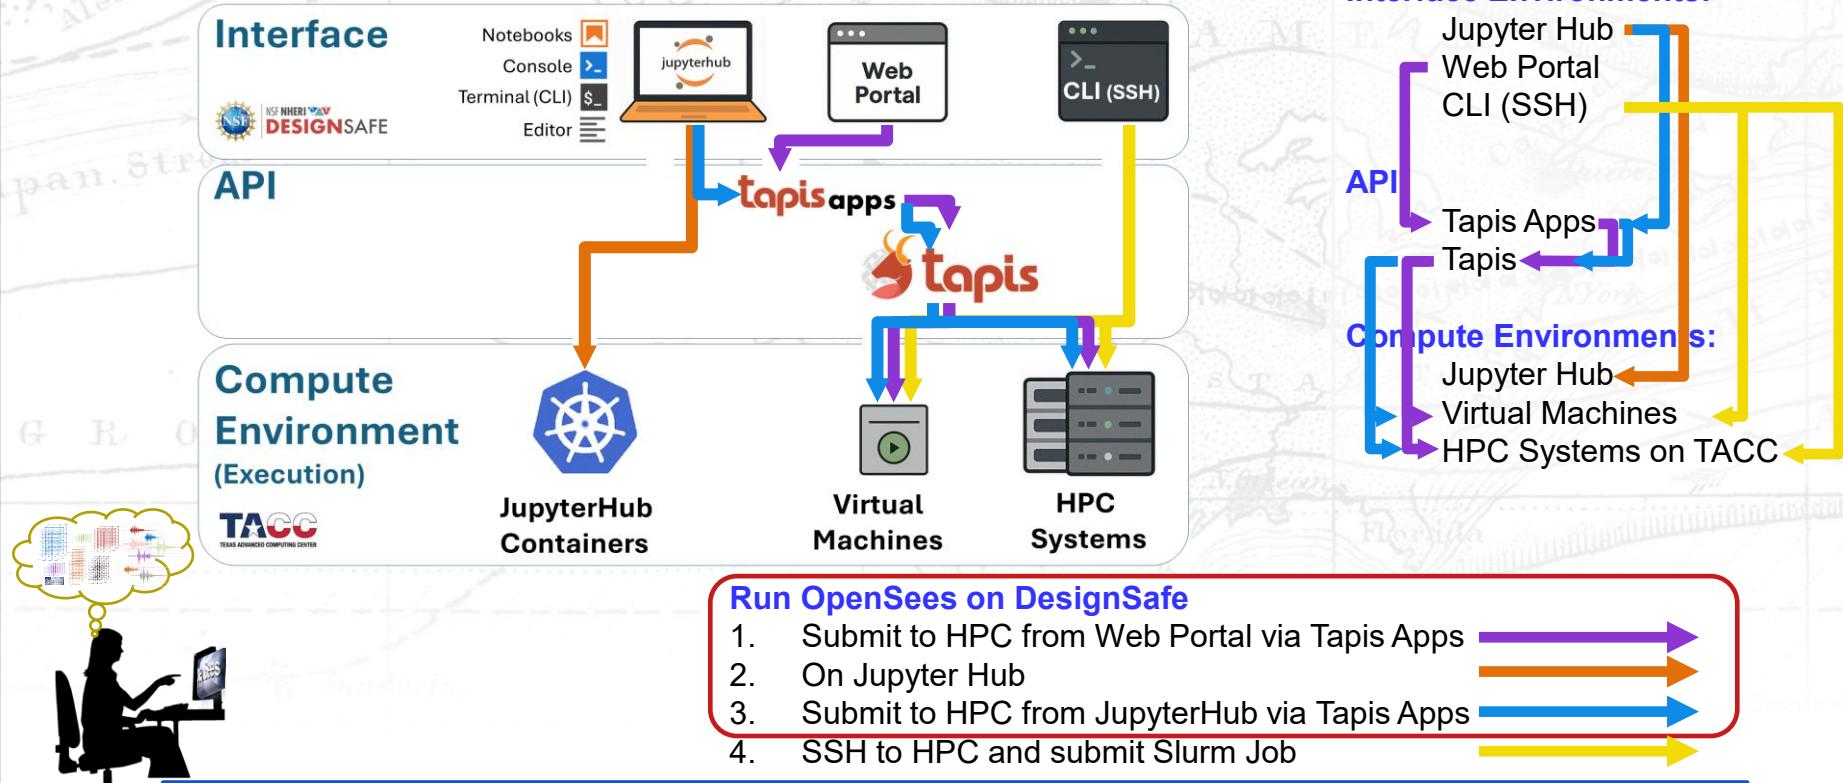
UCLA

TACC

RICE

Florida Tech

Ways To Run OpenSees on DesignSafe



NSF NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

Ways To Run OpenSees on DesignSafe

The screenshot shows a web browser window with the URL <https://designsafe-ci.github.io/training-OpenSees-on-DesignSafe/README.html>. The page title is "OpenSees-on-DesignSafe Training" by Silvia Mazzoni, Applications Specialist @ DesignSafe, dated July 2025. The content discusses DesignSafe as a unified platform for natural hazards engineering, focusing on OpenSees as a featured example for its power and popularity in earthquake engineering. It highlights the workflow lifecycle, from small exploratory runs to automated HPC pipelines. The sidebar includes sections for Training Objectives, Guides (DesignSafe & TACC, OpenSees on DesignSafe, JupyterHub, HPC on TACC, File Storage, Tapis), Training Modules (OpenSees from Web Portal, OpenSees on JupyterHub, OpenSees on HPC), Utilities (OpsUtils()), and Table of Contents. At the bottom, there are links for Disclaimer, Version Info, License, and a "Next" button labeled "Training Objectives".

OpenSees-on-DesignSafe Training

by **Silvia Mazzoni**
Applications Specialist @ DesignSafe
July 2025

As a unified platform for natural hazards engineering, **DesignSafe** supports the full research lifecycle by integrating interactive tools, data management services, and high-performance computing resources. Whether you're developing models, running large-scale simulations, or conducting detailed post-processing, DesignSafe provides the infrastructure to streamline and scale your workflow.

In this training module we will focus on **OpenSees**. Not just because of its power and popularity in earthquake engineering, but because it provides an ideal case study for scalable computational workflows. Built from the ground up for parallel execution and flexible scripting, OpenSees illustrates many common challenges and opportunities faced by researchers in other fields — such as managing complex input files, launching parameter studies, and organizing results across systems.

While OpenSees is the featured example throughout this module, **the real emphasis is on the complete Scientific Workflow lifecycle**. You'll gain hands-on experience with the tools and strategies required to build, manage, submit, monitor, and scale computational research — from small exploratory runs to fully automated HPC pipelines. The goal is to help you understand and orchestrate every step of the workflow, regardless of the specific application.

These foundational concepts are widely transferable: whether you're working with CFD solvers, climate models, structural simulations, or your own custom software, **DesignSafe's workflow infrastructure enables scalable, reproducible, and automated research**.

[Disclaimer](#) >
[Version Info](#) >
[License](#) >

Next >
Training Objectives



Training Resources - DesignSafe | OpenSees-on-DesignSafe Train | +

designsafe-ci.org/user-guide/training/

NSF NHERI DESIGN SAFE

USER GUIDE
designsafe-ci.org

Search docs

DesignSafe Essentials

- Getting Started
- Account Help
- Training Resources
 - DesignSafe FAQ
 - How to Cite DesignSafe

DesignSafe Data Depot

- Overview
- Transferring Your Data
- Steps to Curate and Publish Your Datasets
- Best Practices for Data Curation and Publication
- Data Dissemination and Impact
- Resources for Users
- About the Data Depot
- Policies

Tools and Apps

- Requesting New Applications
 - Popular
 - Simulation
 - Analysis

DesignSafe Essentials > Training Resources

Suggest an update via GitHub

TRAINING RESOURCES

DesignSafe-CI offers a suite of training resources aimed at equipping researchers and engineers with cutting-edge computational skills. These repositories cover a range of topics, including accelerating Python applications, utilizing database APIs, implementing physics-informed neural networks (PINNs), understanding explainable AI (XAI), and applying Deep Operator Networks (DeepONet). Each module provides hands-on tutorials and practical examples to enhance your proficiency in leveraging these technologies for natural hazards research.

OpenSees on DesignSafe

Learn to use OpenSees to develop, run, and scale structural & geotechnical simulations on DesignSafe/TACC. This hands-on module covers managing input files, launching parameter studies, and submitting and monitoring HPC jobs for reproducible workflows.

Database API training

Master database interaction through APIs using DesignSafe's infrastructure. Learn to programmatically access, query, and manipulate research data using RESTful APIs. Includes examples of data retrieval, filtering, and integration with analysis workflows.

Accelerating Python using Cython, Numba and JAX

Optimize your Python code performance using various acceleration techniques. Learn to use Cython for C-level performance, Numba for just-in-time compilation, and JAX for automatic differentiation and parallelization. Includes practical examples and performance comparisons.

Training on XAI using DesignSafe

Explore Explainable AI (XAI) techniques applied to natural hazards engineering. Learn how to interpret machine learning models using methods like LIME, SHAP, and feature importance analysis. The training includes practical examples using Python libraries and real-world datasets.

Physics Informed Machine Learning

Discover Physics-Informed Neural Networks (PINNs) for solving partial differential equations. This training demonstrates how to incorporate physical laws into neural networks, with applications in structural dynamics and fluid mechanics using TensorFlow.

DeepONet Training

Learn about Deep Operator Networks (DeepONet), a novel deep learning framework for learning operators. This training covers the fundamentals of DeepONet architecture, its applications in solving differential equations, and practical implementation using PyTorch.



NSF NHERI DESIGN SAFE



UCLA

TACC

RICE

Florida Tech

The screenshot shows a web browser window displaying the 'Training Objectives' page from the DesignSafe website. The page has a blue header bar with the title 'Training Objectives — OpenSees' and the URL 'designsafe-ci.github.io/training-OpenSees-on-DesignSafe/Docs_MD/View_TrainingObjectives.html'. The main content area features the DesignSafe logo and a search bar. A sidebar on the left contains a navigation menu with sections like 'Training Objectives', 'Guides', 'Training Modules', and 'Utilities'. The main content area has a large heading 'Training Objectives' and several paragraphs of text describing the module's goals and content. A list of bullet points follows, detailing specific learning objectives. At the bottom, there is a section titled 'Training Material'.

Training Objectives

In this training module, you'll learn **how to "run" OpenSees from different computational environments within DesignSafe** — whether that's the Jupyter Hub, the Web Portal, Tapis, or direct HPC command-line access.

Beyond OpenSees-specific workflows, this training covers the **core fundamentals necessary to build scalable and efficient scientific workflows** that optimize the resources available through **TACC and DesignSafe**. We'll discuss how to:

- Understand the **hardware architecture at TACC**, and how to best match job types to system capabilities.
- Use the **"software" infrastructure** that supports computation at scale — including **Tapis**, the **Tapis API**, and **task-specific Tapis Apps** available on DesignSafe.
- Navigate **file storage systems** on DesignSafe, including data staging, transfer, and archival strategies.
- Interact with different computational environments — from terminal-based CLI to Jupyter notebooks and batch script submissions — and know when to use each.

By the end of this module, you'll know how to:

- Use the **DesignSafe Web Portal** to submit OpenSees jobs without writing any batch scripts.
- Launch OpenSees from within a **Jupyter notebook**, passing in variables or dynamically generating input files.
- Run OpenSees directly from a terminal or Jupyter notebook, using both Tcl and Python scripts.
- Submit OpenSees jobs to the HPC system on TACC via **SLURM**, providing input files and specifying resources like cores and wall time.
- Automate large parameter studies by modifying script values on the fly or using loops in Python.
- Leverage **Tapis APIs** to submit and monitor multiple jobs programmatically.

Training Material



NSF NHERI
DESIGNSAFE



TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech


 Search

 Ctrl + K

OpenSees-on-DesignSafe Training

Training Objectives

Guides

[DesignSafe & TACC](#)
[OpenSees on DesignSafe](#)
[JupyterHub](#)
[HPC on TACC](#)
[File Storage](#)
[Tapis](#)

Training Modules

[OpenSees from Web Portal](#)
[OpenSees on JupyterHub](#)
[OpenSees on HPC](#)

Utilities

[OpsUtils\(\)](#)
[Table of Contents](#)

Training Material

Using the Training Notebooks

The pages in the training modules are **generated directly from Jupyter notebooks**. These notebooks are designed to be **templates for your own work** — you should never start from scratch. Instead:

- The notebooks are stored in the **DesignSafe Community Data folder**.
 - You can run them there, but you **cannot save changes** in Community Data.
 - Also, some notebooks include inputs such as job definitions or file paths that are specific to *my user account*.
- To work effectively, you should **copy the notebooks** to your own storage (e.g., `~/MyData`).
- **Edit, duplicate, or modify** the appropriate cells to fit your workflow.

You can recognize a notebook-based page because it has an “**Open in DesignSafe**” button at the very top.

Clicking it will bring you to **JupyterHub in DesignSafe**, where you’ll be prompted to log in (or create an account). I recommend copying **only relevant notebooks in the folder** into your path, as I may be editing the existing notebooks or adding new ones.

Python Function Library

Alongside the notebooks, this training relies on a **custom library of Python functions** that I have developed.

This library is stored in **DesignSafe Community Data** and is used extensively throughout the modules.

- You *can* copy the library directly from Community Data, but I **don't recommend it**, since I am continually updating and improving it.
- Instead, look at the **script near the top of any notebook** that uses the library — copy that snippet into your own path so you can use a stable version of the library in your work.

[Previous
OpenSees-on-DesignSafe Training](#)

[Next
DesignSafe & TACC](#)

By Silvia Mazzoni
© Copyright 2025.

The screenshot shows the DesignSafe Community Data interface. At the top, there's a navigation bar with links for 'Training', 'Using', 'Python', 'Use DesignSafe', 'Learning Center', 'NHERI Facilities', 'NHERI Community', 'News', and 'Help'. Below the navigation is a search bar labeled 'Search Data Files' with a magnifying glass icon. The main area is titled 'Community Data / OpenSees / TrainingMaterial / training-OpenSees-on-DesignSafe'. It lists several items under 'File Name': 'Examples_OpenSees', 'Jupyter_Notebooks', and 'OpsUtils'. A blue arrow points to the 'Community Data' button in the sidebar. Another blue arrow points to the 'Jupyter_Notebooks' item in the list. A third blue arrow points to the 'Help' button at the bottom right of the sidebar.



NSF NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

REFERENCE GUIDES

Everything you don't need to know, but really should!



NSF NHERI
DESIGNSAFE

 TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

The image shows six browser tabs, each displaying a different section of the OpenSees-on-DesignSafe training website. The tabs are arranged horizontally and show the same basic layout with a sidebar on the left and a main content area on the right.

Left Sidebar (Common to all tabs):

- Search bar: Q Search Ctrl + K
- NSF NHERI DESIGNSAFE logo
- OpenSees-on-DesignSafe Training
- Training Objectives
- Guides
- DesignSafe & TACC
 - Workflow Architecture
 - Compute Environments
 - Workflow Decision Guide
- OpenSees on DesignSafe
 - OpenSees-Tcl & OpenSeesPy
 - OpenSees Applications
 - Interpreters & Workflows
 - Decision Matrix
 - Command Structure
 - Executable File
 - Input Script File
 - Command-Line Arguments
 - Parallel Execution
 - Parallel Execution: MPI
 - Parallel Execution: ibrun
 - OpenSeesPy Parallel
 - Execution Guide
- JupyterHub
- HPC on TACC
- File Storage
- Tapis
- Training Modules
 - OpenSees from Web Portal
 - OpenSees on JupyterHub
 - OpenSees on HPC
- Utilities
 - OpsUtils()
- Table of Contents

The image shows six browser tabs, each displaying the same website structure for "OpenSees-on-DesignSafe Training". The tabs are labeled from left to right:

- DesignSafe & TACC — OpenSees-on-DesignSafe
- OpenSees on DesignSafe — OpenSees-on-DesignSafe
- JupyterHub — OpenSees-on-DesignSafe
- HPC on TACC — OpenSees-on-DesignSafe
- File Storage — OpenSees-on-DesignSafe
- Tapis — OpenSees-on-DesignSafe

The website has a header with the NSF NHERI DESIGNSAFE logo. The main content area includes a search bar and a "Guides" section. The "Guides" section is expanded in all tabs, showing the following categories:

- DesignSafe & TACC** (selected in the first tab)
- Workflow Architecture
- Compute Environments
- Workflow Decision Guide
- OpenSees on DesignSafe
- JupyterHub
- HPC on TACC
- File Storage
- Tapis
- Training Modules**
- OpenSees from Web Portal
- OpenSees on JupyterHub
- OpenSees on HPC
- Utilities**
- OpsUtils()
- Table of Contents

The "OpenSees & TACC" section is expanded in the first tab. The "OpenSees on DesignSafe" section is expanded in the second tab. The "JupyterHub" section is expanded in the third tab. The "HPC on TACC" section is expanded in the fourth tab. The "File Storage" section is expanded in the fifth tab. The "Tapis" section is expanded in the sixth tab.

A yellow arrow points to the "HPC on TACC" section in the fourth tab.



NSF NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

OpenSees Tapis Apps — OpenSees on DesignSafe · designsafe-ci.github.io/training-OpenSees-on-DesignSafe/Docs_MD/OpsApps_Overview.html

OpenSees Tapis Apps

The [OpenSeesMP app template](#) provides a working example of how to wrap the parallel version of OpenSees into a Tapis App.

Structure Overview

This template includes:

- **app-definition.json:** Describes the app's metadata and links it to a system (like Frontera or Stampede3). It defines:
 - *deploymentSystem* and *executionSystem*
 - CLI-style *parameters* and *inputs*
 - Output handling rules
- **app-wrapper.sh:** The actual launch script. It:
 - Loads the OpenSees module (or uses a container)
 - Builds the MPI run command (e.g., *ibrun OpenSeesMP*)
 - Moves inputs and manages outputs
- **README.md:** Provides guidance on how to install the app using the Tapis CLI or API.

Example Use Case

Let's say you want to run a parametric analysis using OpenSeesMP on Stampede3. You can:

1. Register the app in your Tapis tenant using the provided *app-definition.json*.
2. Submit a job specifying:

↑ The OpenSees.tcl input script



NSF NHERI
DESIGNSAFE



UCLA

TACC

RICE

Florida Tech



NSF NHERI
DESIGNSAFE

 TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

TRAINING MODULES

3 Ways to Run OpenSees

OpenSees-on-DesignSafe Train... + designsafe-ci.github.io/training-OpenSees-on-DesignSafe/README.html

OpenSees-on-DesignSafe Training

by **Silvia Mazzoni**
Applications Specialist @ DesignSafe
July 2025

As a unified platform for natural hazards engineering, **DesignSafe** supports the full research lifecycle by integrating interactive tools, data management services, and high-performance computing resources. Whether you're developing models, running large-scale simulations, or conducting detailed post-processing, DesignSafe provides the infrastructure to streamline and scale your workflow.

In this training module we will focus on **OpenSees**. Not just because of its power and popularity in earthquake engineering, but because it provides an ideal case study for scalable computational workflows. Built from the ground up for parallel execution and flexible scripting, OpenSees illustrates many common challenges and opportunities faced by researchers in other fields — such as managing complex input files, launching parameter studies, and organizing results across systems.

While OpenSees is the featured example throughout this module, **the real emphasis is on the complete Scientific Workflow lifecycle**. You'll gain hands-on experience with the tools and strategies required to build, manage, submit, monitor, and scale computational research — from small exploratory runs to fully automated HPC pipelines. The goal is to help you understand and orchestrate every step of the workflow, regardless of the specific application.

These foundational concepts are widely transferable: whether you're working with CFD solvers, climate models, structural simulations, or your own custom software, **DesignSafe's workflow infrastructure enables scalable, reproducible, and automated research**.

[Disclaimer](#)



NSF NHERI DESIGNSAFE

TEXAS
The University of Texas at Austin

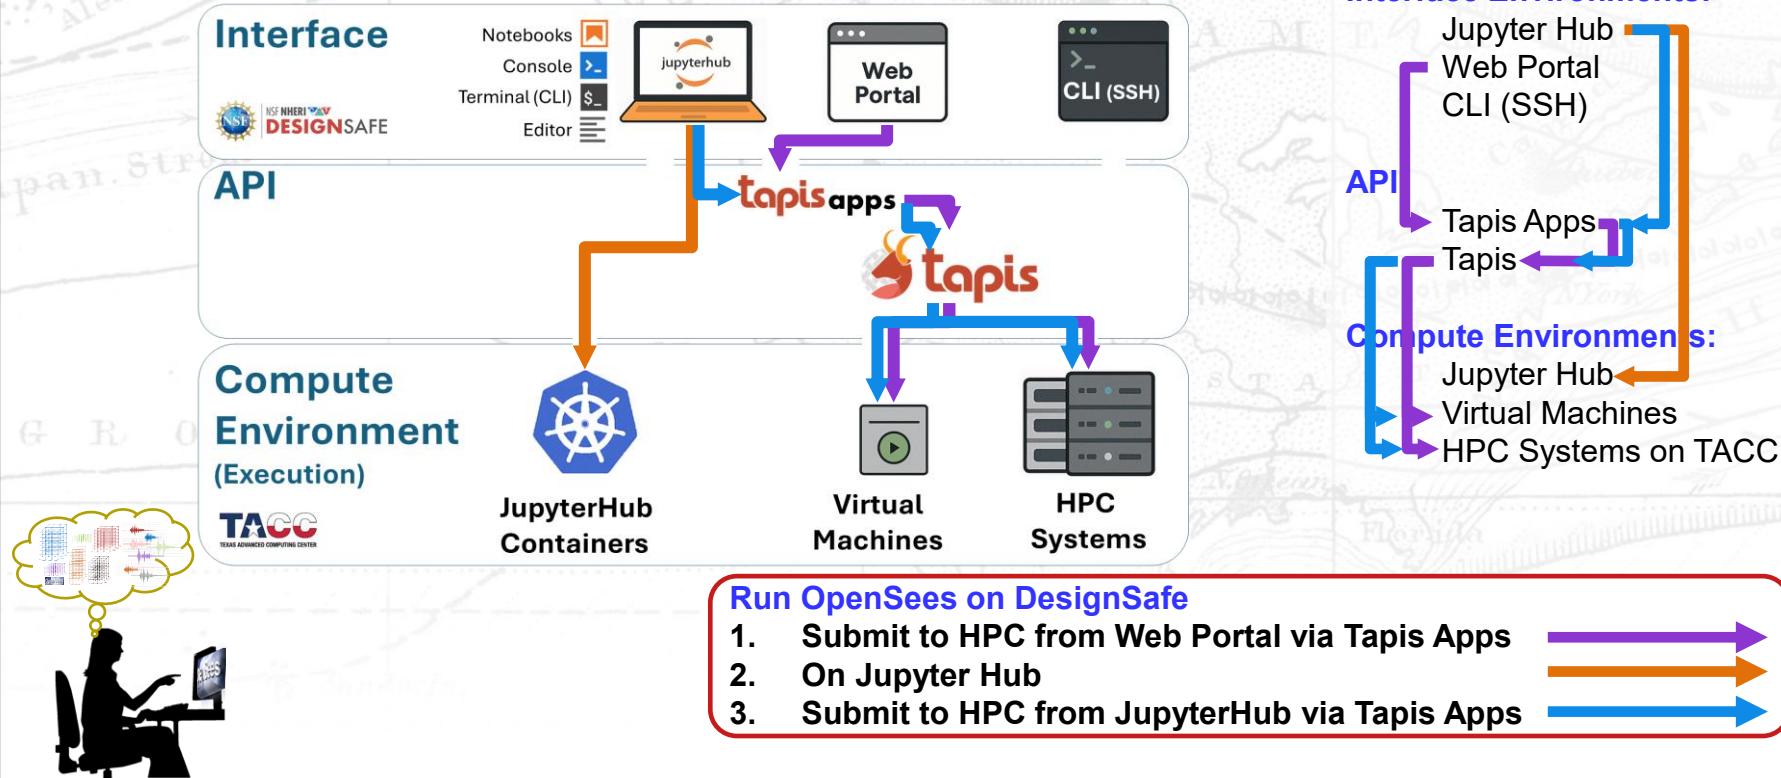
UCLA

TACC

RICE

Florida Tech

Ways To Run OpenSees on DesignSafe



NSF NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

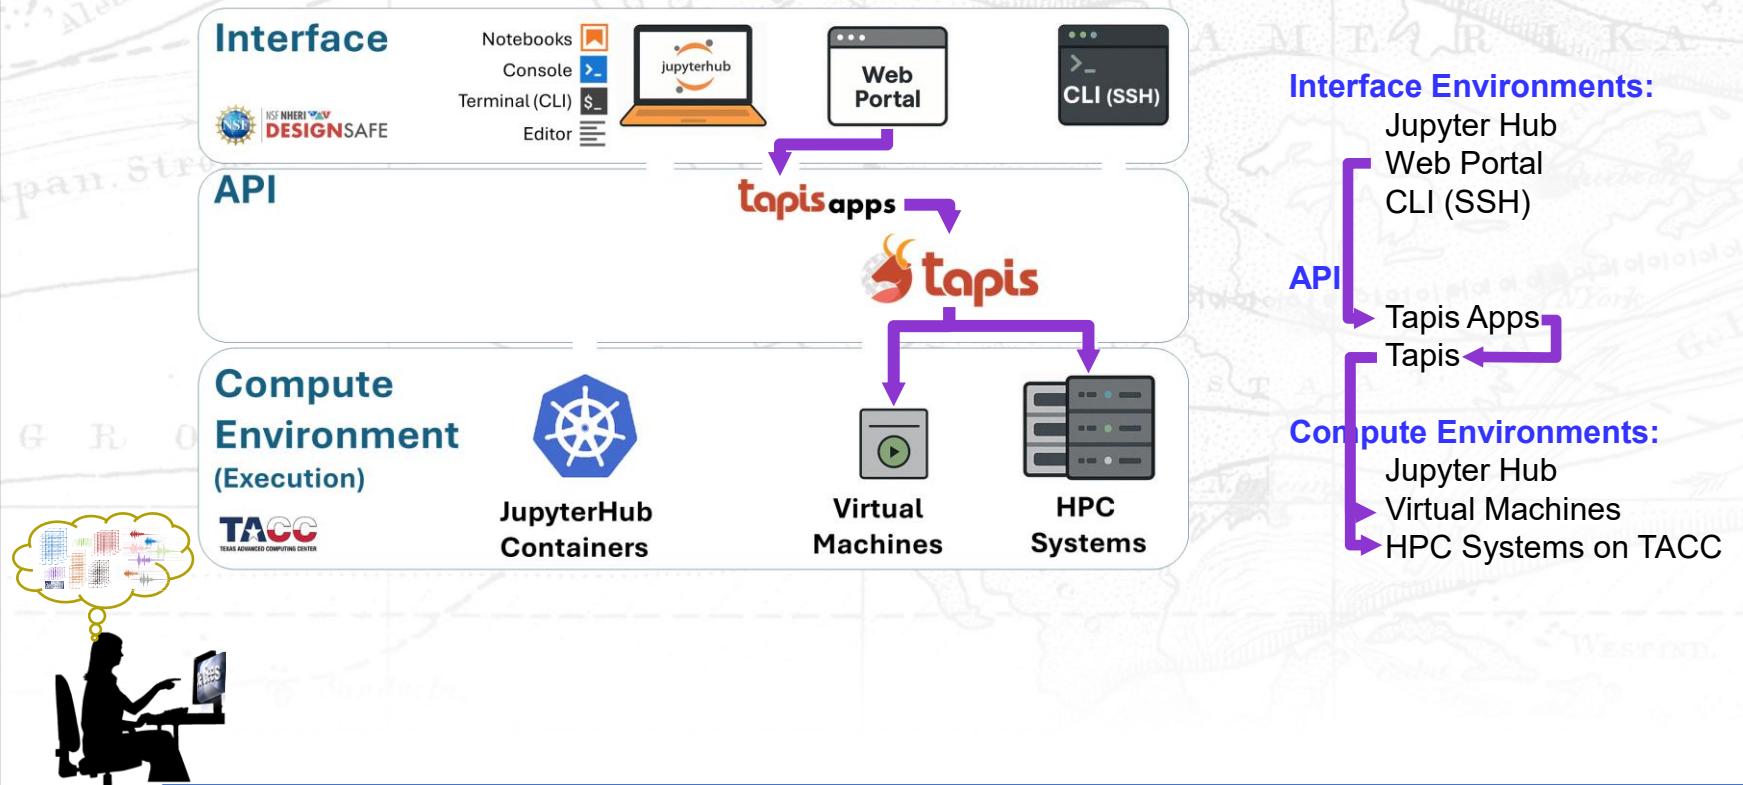
UCLA

TACC

RICE

Florida Tech

1. SUBMIT TO HPC FROM WEB PORTAL

NSF NHERI
DESIGNSAFETEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

OpenSees from Web Portal — designsafe-ci.github.io/training-OpenSees-on-DesignSafe/Docs_MD/WebPortal_Overview.html

OpenSees from Web Portal

Hybrid Workflow with Web Portal + JupyterHub

DesignSafe provides a suite of **Tapis-powered OpenSees applications** accessible through its **Web Portal**, allowing users to run both sequential and parallel jobs on Stampede3 **without writing SLURM scripts manually**. This section introduces the available apps, how they work, and how to integrate them into a **JupyterHub-driven workflow**.

Available Applications

App Name	Description
OpenSees-Express	Sequential execution on a VM; great for small or test models.
OpenSeesMP	Parallel execution using MPI across multiple nodes for large simulations.
OpenSeesSP	Parallel static analysis (domain decomposition).

These apps are backed by the **Tapis platform**, which handles job submission, staging, monitoring, and result retrieval from TACC systems.

Why Use the Web Portal?

The Web Portal simplifies the entire HPC job submission process:

- Upload input files via a browser
- Select the app and main input script
- Configure compute settings (cores, wall time, etc.)



LIVE DEMO



NSF NHERI
DESIGNSAFE

 TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

2. RUN ON JUPYTERHUB

Interface Environments:

Jupyter Hub
Web Portal
CLI (SSH)

API:

Tapis Apps
Tapis

Compute Environments:

Jupyter Hub
Virtual Machines
HPC Systems on TACC

Interface



Web
Portal

CLI (SSH)

API

tapis apps



Compute Environment (Execution)



JupyterHub
Containers



Virtual
Machines



HPC
Systems



NSF NHERI DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

OpenSees on JupyterHub — designsafe-ci.github.io/training-OpenSees-on-DesignSafe/Docs_MD/RunOpsInDS_JupyterHub_intro.html

OpenSees on JupyterHub

The Primary Interface for Scripting, Testing, and Orchestrating Your Workflow

The **DesignSafe JupyterHub** is your command center for running OpenSees in all its forms—whether you’re writing Tcl scripts for the traditional interpreter or Python scripts using OpenSeesPy. It is the **primary portal for scripting, visualization, data management, and HPC job orchestration**.

This environment is ideal for:

- Interactive script development
- Pre- and post-processing of model data
- Exploratory analysis and documentation
- Submitting HPC jobs directly via the Tapis API

Each JupyterHub session runs in a dedicated container on a TACC-managed Kube **CPU cores and 20 GB of RAM**—perfect for medium-scale simulations and full-fe workflows.

All OpenSees have been pre-installed and are available in JupyterHub:

- Tcl: OpenSees, OpenSeesMP, and OpenSeesSP
- Python: OpenSeesPy

Choosing the Right Interface for Your Task

JupyterHub offers a suite of interfaces tailored to different stages of your modelin

Task	Recommended interface
Modeling	JupyterHub
Data Management	Tapis
Execution	JupyterHub
Visualization	JupyterHub
HPC Job Orchestration	Tapis

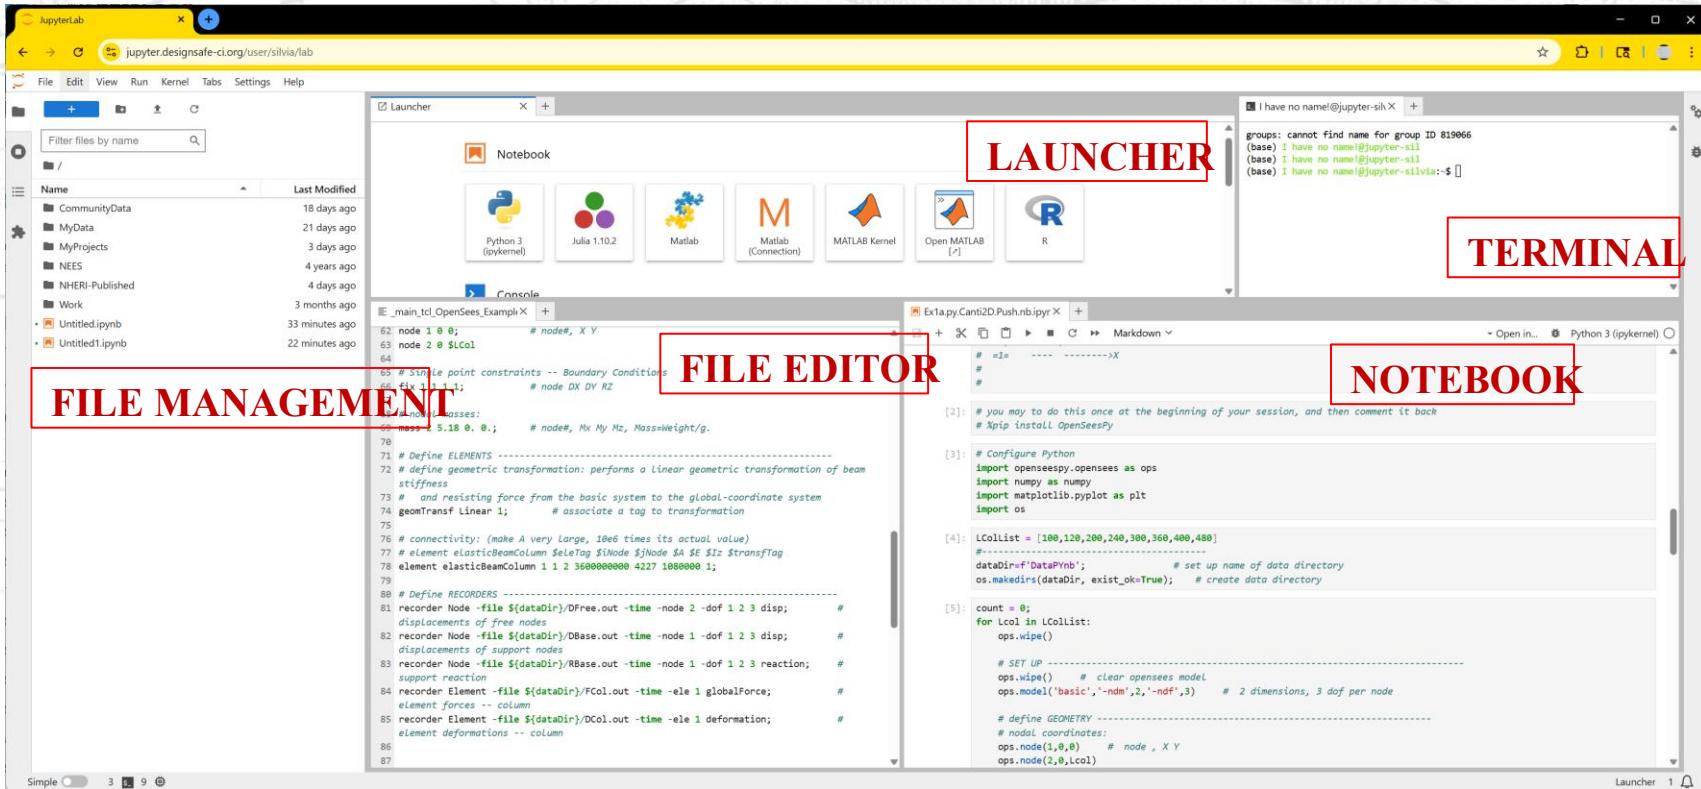
```

graph TD
    subgraph Interface [Interface]
        Notebooks[Notebooks]
        Console[Console]
        Terminal[Terminal]
        Editor[Editor]
    end
    subgraph API [API]
        WebPortal[Web Portal]
        CLISSH[CLI (SSH)]
    end
    subgraph ComputeEnvironment [Compute Environment (Execution)]
        JupyterHubContainers[JupyterHub Containers]
        VMs[Virtual Machines]
        HPC[HPC Systems]
    end
    subgraph Tapis [tapis]
        tapisapps[tapisapps]
        tapis[tapis]
    end

    Notebooks --> tapisapps
    Console --> tapisapps
    Terminal --> tapisapps
    Editor --> tapisapps
    WebPortal --> tapisapps
    CLISSH --> tapisapps
    tapisapps --> tapis
    tapis --> JupyterHubContainers
    tapis --> VMs
    tapis --> HPC
  
```



JUPYTER IS THE BEST! IDE



NSF
NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

OpenSees-on-DesignSafe Training

The screenshot displays a Jupyter Notebook environment with several open tabs and panes:

- File Launcher**: Shows a file tree with categories like CommunityData, MyData, MyProjects, NEES, and NHERI-Published.
- Notebook**: A grid of icons for different kernels: Python 3 (Jupyter), Julia 1.10.2, Matlab, MATLAB (connections), MATLAB Kernel, Open MATLAB, and R.
- Console**: A grid of icons for various kernels: Python 3 (Jupyter), Julia 1.10.2, Matlab, MATLAB (connections), MATLAB Kernel, Open MATLAB, and R.
- Other**: Icons for Terminal, MATLAB File, Text File, Matplotlib File, Julia File, Python File, R File, and Show Contextual Help.
- Terminal**: Shows a command-line session where a user attempts to find a file named 'group_id' but receives an error message: "groups: cannot find name for group ID 0".
- Code Editor**: An OpenSees script titled "Ex01_elasticCantileverPushIn.tcl". The script defines a 3D model of a cantilever column with a fixed base at A and a horizontal force of 2000kip at the free end. It includes definitions for nodes, elements, boundary conditions, and material properties.
- Output**: The results of running the script, showing the calculated displacement and stress distributions along the column.
- Interactive Notebook**: A pane showing the "OpenSees Example 1a: Elastic Cantilever Column - Static Pushover" notebook. It includes an introduction, a diagram of the cantilever column, and a code cell for configuring Python.



NSF NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

LIVE DEMO



NSF NHERI
DESIGNSAFE

 TEXAS
The University of Texas at Austin

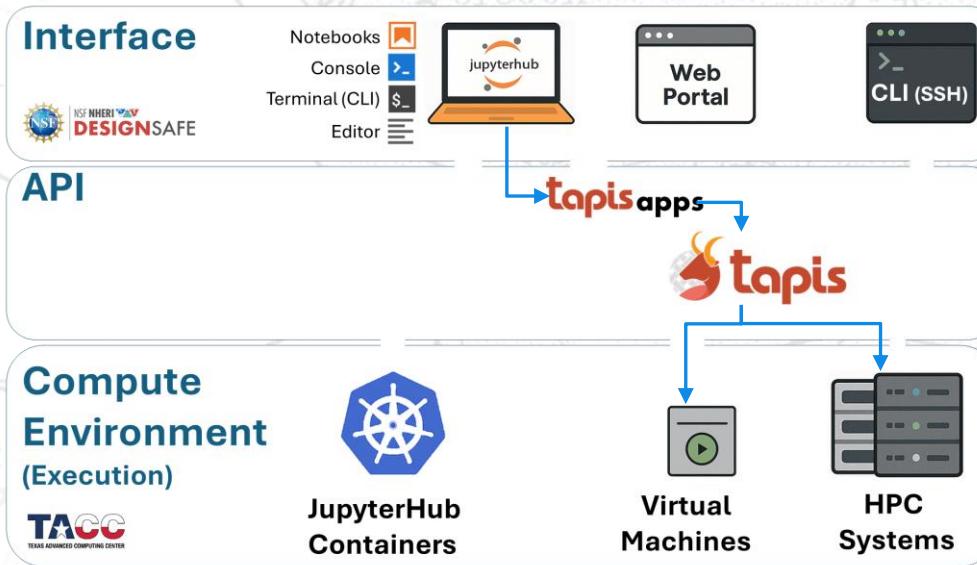
UCLA

TACC

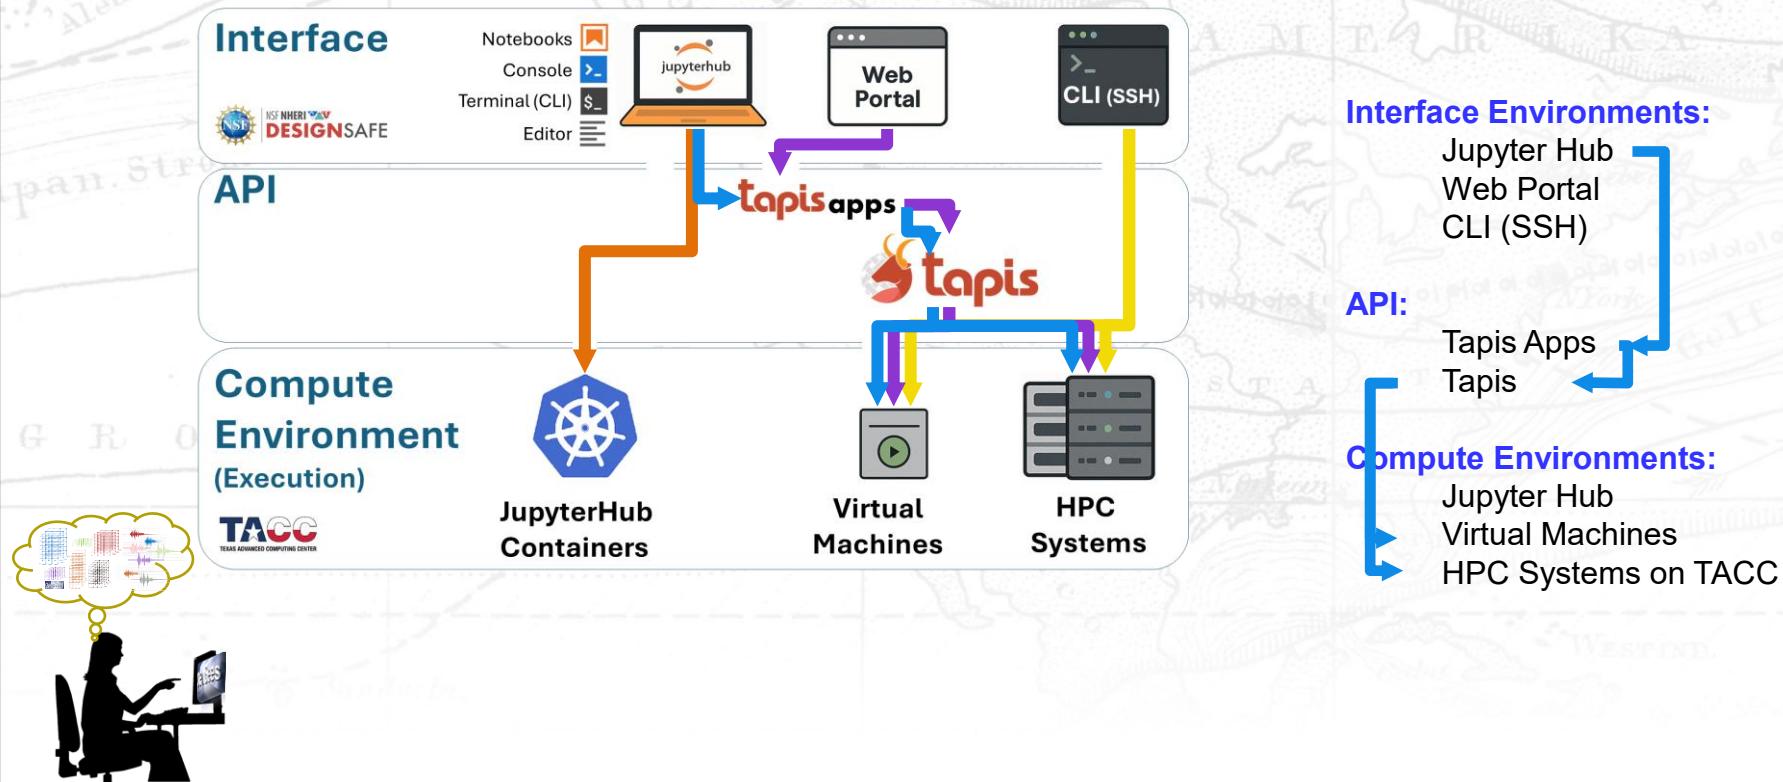
RICE

Florida Tech

3. SUBMIT TO HPC FROM JUPYTERHUB



3. SUBMIT TO HPC FROM JUPYTERHUB



OpenSees on HPC — OpenSees · designsafe-ci.github.io/training-OpenSees-on-DesignSafe/Docs_MD/RunOpenSees_HPCviaJupyterHub.html

OpenSees on HPC

OpenSees on HPC via JupyterHub

DesignSafe's JupyterHub provides an interactive environment where you can prepare, submit, and monitor jobs directly on TACC's HPC systems. Instead of writing your own SLURM scripts, you can use **Tapis Apps**—predefined job templates that take care of loading the correct modules, staging input files, running OpenSees on HPC nodes, and saving results back to your storage.

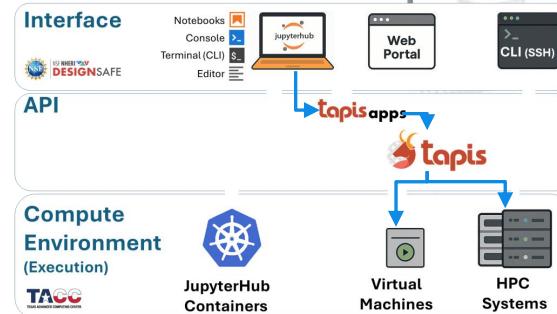
Running OpenSees this way is especially useful when you want to combine the flexibility of Jupyter notebooks with the power of HPC. You can test or pre-process models interactively, then submit larger production runs to Stampede3 without leaving JupyterHub.

When comparing the available workflows:

- **Web Portal → HPC** Easiest to use with a point-and-click interface, but less flexible for automation or scripting.
- **JupyterHub → Local** Good for testing small models interactively, but limited to the resources available in JupyterHub (single node).
- **SSH → HPC (manual SLURM)** The traditional method: log into Stampede3 via SSH, write your own SLURM job script, and submit with `sbatch`. In addition, you must manually move your input files to the correct **scratch location** before submitting the job, and copy the results back afterward. This provides **maximum control**, but requires more setup and becomes tedious when you need to run many jobs.
 - **Note:** Because Tapis makes this process much more fluid, we will not be covering the SSH method in detail in this training.
- **JupyterHub → HPC via Tapis Apps** Combines the best of both worlds: scripting and automation from JupyterHub with the full scalability of HPC. Tapis Apps handle SLURM submission and also take care of staging files in and out of scratch for you, making it especially convenient for large studies or repeated runs.

Step-by-step workflow (JupyterHub + Tapis Apps):

1. Prepare your OpenSees input files in your DesignSafe storage — you can do this within JupyterHub.



Run OpenSees Tapis Apps — designsafe-ci.github.io/training/OpenSees-on-DesignSafe/Docs_MD/tapis/OpenSeesApps.html

Run OpenSees Tapis Apps

OpenSees Tapis Apps on DesignSafe

Running OpenSees on a high-performance computing (HPC) system like Stampede3 isn't as simple as just launching a command. DesignSafe has developed and published three official **OpenSees Tapis Apps** that automate all of this setup and submission:

- **OpenSeesEXPRESS** – for sequential simulations – runs on a dedicated VM
- **OpenSeesMP** – for parallel simulations – runs on Stampede3
- **OpenSeesSP** – for single-processor jobs – runs on Stampede3

The OpenSees app on DesignSafe automates OpenSees Jobs on TACC

The **OpenSees Web-Portal app** on DesignSafe does all of this heavy lifting. It:

- Generates the **SLURM job script for you**, tuned to the TACC environment (whether for *OpenSees*, *OpenSeesMP*, or *OpenSeesSP*).
- Automatically stages your input files to the **HPC scratch directory**, where I/O is fastest.
- Executes your analysis on the compute nodes you requested.
- Collects and returns your output files to your **DesignSafe My Data workspace** after the job completes.

This abstraction allows researchers to launch simulations efficiently—without writing job scripts or accessing the command line on a remote cluster. Whether you're working through the portal, a Jupyter notebook, or a custom workflow, the Tapis apps provide a consistent and scalable way to run OpenSees on TACC resources.

This means you can:

- Focus on your structural or geotechnical model — not on cluster commands.
- Submit jobs through a simple browser interface (or later, programmatically through Tapis or Python).
- Get consistent, optimized performance on TACC hardware without fighting with compilers, environment modules, or manual SLURM scripts.



NSF NHERI
DESIGNSAFE



TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

Query and Retrieve Jobs — OpenSees-on-DesignSafe Training

designsafe-ci.github.io/training-OpenSees-on-DesignSafe/Docs_MD/tapis_query/jobs.html

 **DESIGNSAFE**

Search Ctrl + K

OpenSees-on-DesignSafe Training
Training Objectives
Guides
DesignSafe & TACC
OpenSees on DesignSafe
JupyterHub
HPC on TACC
File Storage
Tapis
Training Modules
OpenSees from Web Portal
OpenSees on JupyterHub
OpenSees on HPC
Set up Tapis
Run OpenSees Tapis Apps
Query and Retrieve Jobs
Step 1: Explore All Jobs
Step 2: Inspect Job
Step 3: Retrieve Output
Explore Jobs Interactively
Cancel Tapis Job
Utilities
OpsUtils()
Table of Contents

Query and Retrieve Jobs

Explore, Inspect, and Download Outputs with the Tapis v3 Python SDK (Tapipy)

Tapis allows you to retrieve both **metadata** and **actual data (outputs)** for all your jobs — whether they were submitted directly through Tapis or via any DesignSafe application (such as the OpenSeesMP and OpenSeesSP web portal apps).

Three levels of job information

When working with Tapis, you typically follow a **three-step process** to explore and extract job data. Each step has dedicated functions suited for that stage:

1. **Explore all jobs (overview metadata)**
 - **List jobs** — get overview metadata across many jobs.
 - Use `getJobList()` to retrieve high-level metadata for **all your jobs**, such as job names, UUIDs, submission times, statuses, and applications used.
 - This is perfect for building a summary DataFrame to explore your jobs and decide which one to examine in detail.
2. **Get detailed metadata for a specific job**
 - **Inspect a specific job** — retrieve full configuration and lifecycle.
 - Once you identify a job of interest (via its UUID), you can retrieve full metadata and execution details with:
 - `getJob(uuid)` for complete job configuration and input/output settings.
 - `getJobStatus(uuid)` for current execution state and progress.
 - `getJobHistory(uuid)` to see a timeline of events and status changes — helpful for performance

Query and Retrieve Jobs — OpenSees-on-DesignSafe · designsafe-ci.github.io/training-OpenSees-on-DesignSafe/Docs_MD/tapis_queryJobs.html

DesignSafe & TACC
OpenSees on DesignSafe
JupyterHub
HPC on TACC
File Storage
Tapis
Training Modules
OpenSees from Web Portal
OpenSees on JupyterHub
OpenSees on HPC
Set up Tapis
Run OpenSees Tapis Apps
Query and Retrieve Jobs
Step 1: Explore All Jobs
Explore All Jobs
Filter Tapis Jobs
Step 2: Inspect Job
Job Status
Job Metadata
Job History
Step 3: Retrieve Output
Access Output Data
List All Job Output
Download All Job Output
Explore Jobs Interactively
Cancel Tapis Job
Utilities
OpsUtils()
Table of Contents

web portal apps).

Three levels of job information

When working with Tapis, you typically follow a **three-step process** to explore and extract job data. Each step has dedicated functions suited for that stage:

- 1. Explore all jobs (overview metadata)**
 - **List jobs** — get overview metadata across many jobs.
 - Use `getJobList()` to retrieve high-level metadata for **all your jobs**, such as job names, UUIDs, submission times, statuses, and applications used.
 - This is perfect for building a summary DataFrame to explore your jobs and decide which one to examine in detail.
- 2. Get detailed metadata for a specific job**
 - **Inspect a specific job** — retrieve full configuration and lifecycle.
 - Once you identify a job of interest (via its UUID), you can retrieve full metadata and execution details with:
 - `getJob(uuid)` for complete job configuration and input/output settings.
 - `getJobStatus(uuid)` for current execution state and progress.
 - `getJobHistory(uuid)` to see a timeline of events and status changes — helpful for performance tuning or debugging failures.
- 3. Access and download job outputs**
 - **Retrieve outputs** — browse and download results.
 - After you've drilled into a specific job, use:
 - `getJobOutputList(uuid)` to list all output files and folders produced by the job.
 - `getJobOutputDownload(uuid, outputPath)` to download individual files to your local workspace



LIVE DEMO



NSF NHERI
DESIGNSAFE

 TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech

USER RESOURCES ON DESIGNSAFE

- User Guides
- Training Modules
- Webinars
- Running-OpenSees-on-DesignSafe Slack Channel
- OpenSees Slack Channel
- Submit a ticket for issues
- Message me on slack. (I prefer the channel so others can participate)



NSF NHERI
DESIGNSAFE



UCLA

TACC

RICE

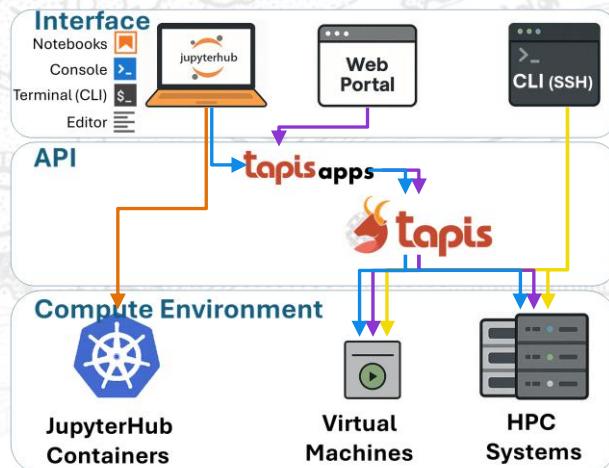
Florida Tech

SUMMARY & QUESTIONS

DesignSafe provides the computational environment to integrate your OpenSees analyses into your complete workflow

Recommended Getting Started:

1. With a MWE – minimum working example
2. Test the MWE in Jupyter Hub
3. Personalize and test your script in JupyterHub
4. Submit your modified script to the Web Portal
5. Use Tapis to query your Job Metadata
6. Access your analysis results in Jupyter Hub
7. Submit the SAME script to HPC via Tapis



NSF NHERI
DESIGNSAFE

TEXAS
The University of Texas at Austin

UCLA

TACC

RICE

Florida Tech